



IBM Information Management software

Locks and concurrency




Why are locks needed?

- Ensure data integrity while permitting multiple applications to access data
- Prohibit applications from accessing uncommitted data written by other applications (*unless* Uncommitted Read isolation level is used)
- Control undesirable effects
 - Lost data due to concurrent updates
 - Unrepeatable reads
 - Phantom read phenomenon



Locking strategies



DB2 employs *either* strict table locking **OR** table locking in conjunction with row locking for typical application processing

Table lock modes

| | |
|-----|-----------------------------|
| IN | Intent None |
| IS | Intent Share |
| IX | Intent eXclusive |
| SIX | Share with Intent eXclusive |
| S | Share |
| U | Update |
| X | eXclusive |
| Z | superexclusive |

Row Locking also used

Strict Table Locking

(See next page)

Row lock modes

| Row Lock | | Minimum* Supporting Table Lock |
|-----------|-------------------------|--------------------------------|
| S | Share | IS |
| U | Update | IX |
| X | eXclusive | IX |
| W | Weak exclusive | IX |
| NS | Next key Share | IS |
| NW | Next key Weak exclusive | IX |

An application does not acquire
if it is using Table Locks of

Row locks

S, U, X, or Z

Lock mode compatibility

| MODE OF LOCK A | MODE OF LOCK B | | | | | | | |
|----------------|----------------|-----|-----|-----|-----|-----|-----|----|
| | IN | IS | S | IX | SIX | U | X | Z |
| IN | YES | YES | YES | YES | YES | YES | YES | NO |
| IS | YES | YES | YES | YES | YES | YES | NO | NO |
| S | YES | YES | YES | NO | NO | YES | NO | NO |
| IX | YES | YES | NO | YES | NO | NO | NO | NO |
| SIX | YES | YES | NO | NO | NO | NO | NO | NO |
| U | YES | YES | YES | NO | NO | NO | NO | NO |
| X | YES | NO | NO | NO | NO | NO | NO | NO |
| Z | NO | NO | NO | NO | NO | NO | NO | NO |

Table Locks

Row Locks

| LOCK A MODE | MODE OF LOCK B | | | | | |
|-------------|----------------|-----|----|-----|-----|-----|
| | S | U | X | W | NS | NW |
| S | YES | YES | NO | NO | YES | NO |
| U | YES | NO | NO | NO | YES | NO |
| X | NO | NO | NO | NO | NO | NO |
| W | NO | NO | NO | NO | NO | YES |
| NS | YES | YES | NO | NO | YES | YES |
| NW | NO | NO | NO | YES | YES | NO |




Introduction to isolation levels



- DB2 provides different levels of protection to isolate data:
 - Uncommitted read (UR)
 - Cursor stability (CS) - default
 - Read stability (RS)
 - Repeatable read (RR)
- Cursor stability is the default isolation level
- Isolation level can be specified for a session, a client connection, or an application before a database connection
 - For embedded SQL, the level is set at bind time
 - For dynamic SQL, the level is set at run time

Isolation levels

| ID | NAME | QTY |
|----|-------|-----|
| 1 | DISP | 10 |
| 2 | KEYB | 3 |
| 3 | MOUSE | 15 |
| 4 | CABLE | 18 |
| 5 | CPU | 1 |
| 6 | SOUND | 4 |



NO ROW
LOCKS
FOR
READ-ONLY
CURSORS

COMMIT POINT

Application Bound



RS or RR

CS

UR

W or X-ROW LOCKS ALWAYS HELD UNTIL COMMIT

DB2 isolation levels



1. UR - Uncommitted Read

- For read only queries, no record locking
- Will see uncommitted changes by other transactions
- Good for accessing read only tables
- Statements in UR which modify data are upgraded internally to CS

2. CS - Cursor Stability

- Default isolation level
- Locks and unlocks each row, 1 at a time (never has 2 locks at once)
- Guaranteed to only return data which was committed at the time of the read

3. RS - Read Stability

- Will keep all qualifying rows locked until the transaction is completed
- Does release locks on rows that do not satisfy query predicates
- Use for result set stability or when future actions on returned rows may be taken

4. RR - Repeatable Read

- Will lock all rows visited and keep locks until the transaction is completed
- Use only when consistent results are required

DB2 and ANSI isolation levels

anomalies allowed and disallowed

| DB2 Isolation | ANSI Isolation | Dirty Write | Dirty Read | Fuzzy Read | Phantom Read |
|-----------------------|----------------------------|-------------|------------|------------|--------------|
| Uncommitted Read (UR) | Read Uncommitted (Level 0) | × | ✓ | ✓ | ✓ |
| Cursor Stability (CS) | Read Committed (Level 1) | × | × | ✓ | ✓ |
| Read Stability (RS) | Repeatable Read (Level 2) | × | × | × | ✓ |
| Repeatable Read (RR) | Serializable (Level 3) | × | × | × | × |

Isolation levels and Read locks for Query results

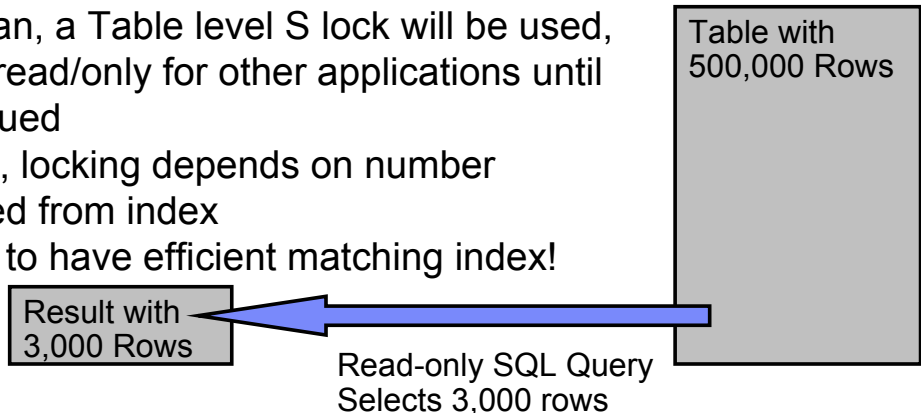
- Uncommitted Read – 1 lock needed IN Table lock, No Row locks
- Cursor Stability – 2 Locks needed
IS Table lock, 1 NS row lock
- Read Stability – 3001 locks needed
IS Table lock, 3000 NS Row Locks

- Repeatable Read

For a Table scan, a Table level S lock will be used, so the table is read/only for other applications until a commit is issued

For Index scan, locking depends on number of rows matched from index

Very important to have efficient matching index!




SET CURRENT ISOLATION

- Assigns a value to the CURRENT ISOLATION special register
`SET [CURRENT] ISOLATION [=] {UR | CS | RR | RS | RESET}`
- CURRENT ISOLATION can be set anytime while connected to a database (special registers are attributes of the connection session)
- Applies to the current session only
- The register setting is used immediately for subsequent SQL statements *until value is set again*, thus allowing different isolation to be applied to different SELECT (and other) statements even within one UOW/transaction
- You can read the current register setting with:

`VALUES CURRENT ISOLATION`

Lock wait and timeout




If the application *hogging* the locks
doesn't COMMIT or ROLLBACK
other applications wait until lock is
available or timeout exceeded

Deadlock causes and detection

UNIT OF WORK —

DELETE SOME CEREAL AND MILK



dchktime

10000

milliseconds

Deadlocks

- Deadlock detector wakes up every DLCHKTIME
 - Each database partition has its own deadlock detector (*db2dlock* process)
 - Catalog database partition (DPF) has the global deadlock detector (*db2glock*)
- A deadlock occurs when there is a cycle between 2 or more transactions which are in lock wait on one another - and the deadlock detector is able to wake up and recognize this before any lock timeout occurs
- DB2 itself will determine which transaction among those involved in the deadlock cycle will be chosen to be the deadlock victim and rolled back
 - Some utilities are executed with an internal flag, to **not** be chosen as the victim
 - Preferential treatment is given to transactions holding Z locks
 - There is currently no way to *help* DB2 choose or prefer/save a victim
- The victim will rollback and return SQL error SQL0911 or SQL0913
- A Deadlock Event monitor can be used to capture information when deadlocks are detected.
- The Health Monitor tracks deadlocks per hour and can generate an alert if the deadlock rate is unusually high.

Snapshot Monitor — taking snapshots

